

SPP Assignment 2

2023101045

Note: Throughout the assignment, only the official `detect.py` was used to make any inference, and it was modified for optimization.

System Specs

I chose to use google colab for bench marking the models as my laptop lacked a dedicated GPU.

Specs:

- CPU : Intel(R) Xeon(R) CPU @ 2.00GHz
- GPU : Tesla T4
- RAM: 12.7 GB

Data Set Used

I used 100 images from the COCO dataset.

Link: <http://images.cocodataset.org/zips/val2017.zip>

Latency - FPS Output

Performance comparison:

model	latency	fps
0 yolov5n	14.1	7.993569
1 yolov5s	15.0	9.390652
2 yolov5m	23.7	8.512493
3 yolov5l	25.7	8.265063
4 yolov5x	41.2	6.480341

GFLOPS Calculation

```
# Print the results
print("The GPU used has a peak flop performance of 8.1 TFLOPS/s - (T4 GPU on google colab)")
print(model_stats_df)
```

The GPU used has a peak flop performance of 8.1 TFLOPS/s - (T4 GPU on google colab)

	Model	Parameters (M)	Model Size (MB)	GFLOPS per inference	GFLOPS/sec \
0	yolov5n	1.87	7.12	4.50	318.87
1	yolov5s	7.23	27.56	16.49	1099.01
2	yolov5m	21.17	80.77	48.97	2066.13
3	yolov5l	46.53	177.51	109.15	4246.89
4	yolov5x	86.71	330.75	205.67	4991.97

	Utilization (%)	Bound Type
0	0.04	memory
1	0.14	memory
2	0.26	memory
3	0.52	memory
4	0.62	memory

Clearly, the maximum utilization is not achieved and hence the inference process is memory bound.

I have done the per layer granularity evaluation for yolov5s model, and here are my results:

The layers with highest utilisation are:

- *model.model.20.m.0.cv2.conv (42.92%)*
 - This convolutional layer has the highest utilization at nearly 43%
 - It achieves 128.75 GFLOPS/s performance.
 - This indicates efficient hardware usage for this particular layer
- *model.model.4.m.1.cv2.conv (40.83%)*
 - Close second with ~41% utilization
 - Shows strong compute efficiency with 122.50 GFLOPS/s
- *model.model.8.m.0.cv2.conv (40.31%)*
 - Maintains high efficiency at ~40% utilization
 - Performs at 120.92 GFLOPS/s

The layers with least utilisation are:

- *model.model.11 (0.16%)*

- Extremely low utilization of just 0.16%
- Very poor performance at only 0.47 GFLOPS/s
- And many more...

Several layers show 0.00% utilization, particularly activation layers (with .act suffix) which are likely implemented using operations that don't require significant compute and might be fused with other operations.

Performance Insights:

- Convolutional layers show high compute utilization (30-43%), while activation layers are memory-bound with negligible utilization.
- **Bottlenecks:** Overall model utilization is just 16.85%, indicating significant optimization potential.
- Memory bandwidth is likely the limiting factor, with uneven workload distribution (utilization ranging from 0.16% to 42.92%).
- Throughput: The current inference time of 162.61ms gives 6.15 FPS, which is far below the potential of the hardware. Average performance is only 50.54 GFLOPS/s, suggesting room for better utilization.

Task 3

For yolov5s

1. Torch.Profiler

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total	CUDA time avg	CPU Mem	Self CPU Mem	CUDA Mem	Self CUDA Mem	# of Calls
model_inference	35.39%	6.488ms	99.91%	18.315ms	18.315ms	0.000us	0.00%	11.833ms	11.833ms	0 b	0 b	92.50 Kb	179.55 Mb	1
aten::conv2d	1.37%	251.790us	48.60%	8.980ms	148.466us	0.000us	0.00%	10.085ms	166.754us	0 b	0 b	98.47 Mb	0 b	60
aten::convolution	1.28%	235.863us	47.22%	8.656ms	144.278us	0.000us	0.00%	10.085ms	166.754us	0 b	0 b	98.47 Mb	0 b	60
aten::convolution	4.44%	814.420us	45.94%	8.421ms	140.352us	0.000us	0.00%	10.085ms	166.754us	0 b	0 b	98.47 Mb	0 b	60
aten::cudnn_convolution	26.49%	4.856ms	33.18%	6.082ms	161.363us	8.891ms	75.14%	8.891ms	148.186us	0 b	0 b	98.47 Mb	98.47 Mb	60
cudnnLaunchKernel	11.88%	2.174ms	11.90%	2.174ms	7.764us	0.000us	0.00%	0.000us	0.000us	0 b	0 b	0 b	0 b	280
aten::add	4.85%	741.965us	6.29%	1.152ms	19.283us	1.114ms	9.41%	1.114ms	18.568us	0 b	0 b	0 b	0 b	60
aten::silu	2.25%	411.647us	4.40%	887.296us	14.162us	647.831us	5.47%	647.831us	11.365us	0 b	0 b	0 b	0 b	57
aten::cat	2.37%	434.889us	3.76%	698.157us	34.508us	531.226us	4.49%	531.226us	26.561us	0 b	0 b	48.44 Mb	48.44 Mb	20
aten::reshape	8.63%	115.992us	2.03%	372.881us	6.213us	0.000us	0.00%	0.000us	0.000us	0 b	0 b	0 b	0 b	60
aten::arange	8.55%	161.088us	1.88%	345.072us	28.756us	14.588us	0.13%	29.769us	2.488us	0 b	0 b	6.80 Kb	0 b	12
aten::mul	0.99%	182.326us	1.56%	286.109us	19.874us	72.511us	0.61%	72.511us	4.834us	0 b	0 b	789.50 Kb	789.50 Kb	15
aten::view	1.49%	273.628us	1.40%	273.628us	3.648us	0.000us	0.00%	0.000us	0.000us	0 b	0 b	0 b	0 b	75
aten::add	0.57%	105.038us	0.90%	165.421us	16.542us	99.981us	0.84%	99.981us	9.980us	0 b	0 b	9.18 Mb	9.18 Mb	10
cudnnEventRecord	0.99%	165.876us	0.90%	165.076us	2.731us	0.000us	0.00%	0.000us	0.000us	0 b	0 b	0 b	0 b	60
aten::stack	0.10%	18.148us	0.77%	148.436us	46.812us	0.000us	0.00%	19.897us	6.336us	0 b	0 b	66.00 Kb	0 b	3
aten::contiguous	0.04%	8.851us	0.73%	133.702us	44.567us	0.000us	0.00%	208.781us	69.567us	0 b	0 b	8.18 Mb	0 b	3
aten::clone	0.14%	25.283us	0.60%	125.651us	41.884us	0.000us	0.00%	208.781us	69.567us	0 b	0 b	8.18 Mb	0 b	3
cudaMemsetAsync	0.52%	95.929us	0.52%	95.929us	5.643us	0.000us	0.00%	0.000us	0.000us	0 b	0 b	0 b	0 b	17
aten::max_pool2d	0.87%	13.216us	0.47%	88.427us	20.809us	0.000us	0.00%	112.287us	37.429us	0 b	0 b	1.17 Mb	-2.34 Mb	3
aten::select	0.38%	70.112us	0.45%	82.944us	5.530us	0.000us	0.00%	0.000us	0.000us	0 b	0 b	0 b	0 b	15
aten::sub	0.39%	54.354us	0.43%	78.823us	26.088us	13.280us	0.11%	13.280us	4.427us	0 b	0 b	197.00 Kb	197.00 Kb	3
aten::max_pool2d_with_indices	0.26%	46.814us	0.40%	71.211us	24.484us	112.287us	0.93%	112.287us	37.429us	0 b	0 b	3.52 Mb	3.52 Mb	3
aten::pow	0.26%	47.461us	0.39%	71.328us	23.776us	9.312us	0.08%	9.312us	3.184us	0 b	0 b	197.00 Kb	197.00 Kb	3
aten::upsample_nearest2d	0.21%	38.989us	0.37%	68.018us	34.005us	49.824us	0.42%	49.824us	24.912us	0 b	0 b	4.69 Mb	4.69 Mb	2

Self CPU time total: 18.315ms
Self CUDA time total: 11.833ms

Percent of time	Line	Stage
	max_det=max_det)	
Rest		Post-Processing (Printing Results, saving to file etc)

```

for path, im, im0s, vid_cap, s in dataset:
    with dt[0]:
        im = torch.from_numpy(im).to(model.device)
        im = im.half() if model.fp16 else im.float() # uint8 to fp16/32
        im /= 255 # 0 - 255 to 0.0 - 1.0
        if len(im.shape) == 3:
            im = im[None] # expand for batch dim
        if model.xml and im.shape[0] > 1:
            ims = torch.chunk(im, im.shape[0], 0)

```

Most of the time was spent in loading the model. Therefore, `detect.py` was modified to print the times after the model loaded and the following result was obtained:

```

--- Pipeline Breakdown (after model load) ---
Pre-processing:  0.20%
Inference:      14.51%
Post-processing: 85.29%
_ _ _ _ _

```

CPU vs GPU Time

From the output of the torch.profiler, we see that the result is as follows for a particular run on a random test image:

```

Self CPU time total: 17.536ms
Self CUDA time total: 11.835ms

```

Memory or threading issues

TLDR: For YOLOV5s , no memory issues observed.

→ We used `torch.utils.bottleneck` to profile the entire inference pipeline on `yolov5s` (For the full report, see above under task 3)

→ If there were threading bottlenecks (e.g., excessive time in Python's `threading` module or contention between threads), it would be reported in the cProfile section — it's not.

→ The tool did not report any threading bottlenecks or contention, suggesting the current pipeline does not suffer from Python GIL limitations or data loader inefficiencies.

→ Autograd and CUDA profiler outputs confirmed that major operations are dominated by `aten::uniform_` and convolution layers, with no evidence of memory leaks or synchronization stalls.

→ CUDA uses 1.5GB out of 15.36 GB of GPU memory available.

NVIDIA-SMI 550.54.15			Driver Version: 550.54.15			CUDA Version: 12.4		
GPU	Name	Perf	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC
Fan	Temp		Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute	M. MIG M.
0	Tesla T4		Off	00000000:00:04.0	Off			0
N/A	75C	P0	32W / 70W	1526MiB / 15360MiB		0%	Default	N/A

Processes:							GPU Memory
GPU	GI	CI	PID	Type	Process name		Usage
	ID	ID					

System RAM
2.2 / 12.7 GB



GPU RAM
1.5 / 15.0 GB



Disk
41.4 / 112.6 GB



Task 4

Approach 1: FP16

Mixed precision (FP16) inference is a well-known optimization technique that reduces model size and computational cost by using 16-bit floating point numbers instead of the standard 32-bit format. The key expectation is improved throughput (FPS) and reduced latency due to faster memory transfers and increased parallelism, especially on GPUs with Tensor Cores like the NVIDIA T4.

Analysis

Changing to FP16 definitely reduces the size of the model; however the expected performance gain by reduced computation time is not observed **for yolov5s**.

```
[35] # FP32 s  
      print(benchmark('yolov5s'))
```

```
➞ {'model': 'yolov5s', 'latency': 14.8, 'fps': 10.452216485549163}
```

```
▶ # Modify Source code to use FP16
```

```
#FP16 s  
print(benchmark('yolov5s'))
```

```
➞ {'model': 'yolov5s', 'latency': 18.6, 'fps': 9.639478825594844}
```

A potential cause might be that YOLOv5s is already small and fast — FP16 might not bring noticeable benefit. In fact, the worse performance could be due to the overhead that conversion to FP16 brings.

Hence we try for larger models and indeed we see a **massive latency improvement of ~ 50%**. In fact, the latency here also includes the time for loading the model at the beginning. Therefore, the performance gain is slightly better than 50%.

```
[52] #FP32 X
      print(benchmark('yolov5x'))
```

```
⇒ {'model': 'yolov5x', 'latency': 42.8, 'fps': 7.2051700580333184}
```

```
▶ # fp16 X
   print(benchmark('yolov5x'))
```

```
⇒ {'model': 'yolov5x', 'latency': 25.6, 'fps': 7.952850937818522}
```

CONCLUSION: FP16 shines with large batches and complex models.

Model	Latency with FP16 (ms)	Latency with FP32 (ms)
yolov5s	18.6	14.8
yolov5x	25.6	42.8

Approach 2: Batch Processing

Component	Approx. Memory Usage
YOLOv5s model	~28MB
One 640×640 FP32 image	~4.7MB
Intermediate tensors	Depends on batch
Overhead (CUDA, others)	~1.5–2GB

So with ~14GB available after overhead, a **safe batch size estimate**:

- **FP32:** $\text{batch_size} \approx 14,000 / 4.7 \approx 297$ → but realistically, 32–64 is safe due to tensor growth.

For the smaller model, batching did lead to an improvement both in latency and to some extent even for throughput in some runs **however this improvement was not reproducible; the subsequent runs did not produce improvement. Similarly, no improvement was observed for larger models.**


```
[65] # yolov5s Batch size: 1
      print(benchmark('yolov5s'))
```

```
↔ {'model': 'yolov5s', 'latency': 24.2, 'fps': 8.98612761578806}
```

```
▶ # Modify source code
  # yolov5s Batch Size: 32
  print(benchmark('yolov5s'))
```

```
↔ {'model': 'yolov5s', 'latency': 19.3, 'fps': 9.585152053050635}
```

```
[75] # yolov5s Batch size 100
      print(benchmark('yolov5s'))
```

```
↔ {'model': 'yolov5s', 'latency': 15.1, 'fps': 10.759021666929808}
```

```
▶ # yolov5x Batch size 1
  print(benchmark('yolov5x.pt'))
```

```
↔ {'model': 'yolov5x.pt', 'latency': 42.0, 'fps': 7.179382335781906}
```

```
[26] # yolov5x Batch size 32
      print(benchmark('yolov5x.pt'))
```

```
↔ {'model': 'yolov5x.pt', 'latency': 41.5, 'fps': 7.255651637438981}
```

```
[ ] # yolov5x Batch size 100
     print(benchmark('yolov5x'))
```

```
↔ {'model': 'yolov5x', 'latency': 44.0, 'fps': 7.343453094493132}
```

The potential reason is that the official script used for running the inference `detect.py` in YOLOv5 **loads and infers images one-by-one**, even if you point it to a folder with 100 images.

Therefore, I wrote a custom script to completely bypass the `detect.py`. (Note that this does not include model load times like `detect.py`).

The goal of batching is to perform inference on multiple images simultaneously in a single forward pass through the neural network. The code below implements

this by manually constructing a batch tensor and passing it to the YOLOv5 model in one go:

```
batch = torch.stack(imgs).to('cuda').half()
output = model(batch)
```

The performance still did not improve. The results are as follows:

model	Batch Size	Pruned Latency
s	1	8.96
s	16	7.75
s	32	7.85
s	64	8.2
s	100	8.12

Batching is most effective when dealing with large workloads where the computational cost of inference dominates. In our case, we processed a relatively small dataset (e.g., 128 images), and used a model (`yolov5s`) that is already highly optimized for fast, per-image inference. As a result, the total inference time is already quite low, and the additional gains from batching are limited or hidden under other forms of overhead.

Additionally, batching introduces some fixed overhead — such as memory allocation, GPU scheduling, and setup time — which only becomes worthwhile when the batch size is large enough to amortize these costs. For small to moderate batch sizes (e.g., 4 or 8), this overhead may cancel out the theoretical speedup from parallel processing.

In summary: **Batching helps more at scale.**

Approach 3: Using ONNX

Deploying YOLOv5 models using the ONNX format with ONNX Runtime can yield better inference performance compared to running models directly in PyTorch.

This is due to several key factors:

1. **Graph-Level Optimizations:** ONNX Runtime performs optimizations such as operator fusion and constant folding, reducing computational overhead and improving execution speed.
2. **Hardware Acceleration:** ONNX supports execution providers like TensorRT and CUDA, which leverage highly optimized, low-level libraries for GPU inference, often outperforming PyTorch's general-purpose engine.
3. **Reduced Overhead:** ONNX models are statically defined and require less runtime graph construction, resulting in faster initialization and lower latency, particularly for repeated inferences.
4. **Efficient Batch Processing:** ONNX Runtime scales better with larger batch sizes, improving throughput in batch inference scenarios.

In summary, ONNX enables faster and more efficient inference by optimizing the execution graph, reducing overhead, and leveraging hardware-specific acceleration, making it well-suited for high-performance deployment.

Results

Without onnx ~ 40ms

With onnx ~ 24 ms.. The results can be observed on the notebook.